

End-to-End Data Engineering Project: Building a Complete Analytics Pipeline with PySpark and DBT

Mohd Laraib Khan | khan.laraib.laraib77@gmail.com | www.linkedin.com/in/khanofdataengineering

Project Overview

This project demonstrates the construction of a complete data engineering pipeline using a pseudo-Uber dataset. The pipeline follows the Medallion Architecture (Bronze for raw data, Silver for cleaned and transformed data, and Gold for modeled analytics-ready data). It handles incremental data ingestion, data cleaning, deduplication, upserts, and dimensional modeling with slowly changing dimensions (SCD Type 2). The dataset includes CSV files for entities like trips, customers, drivers, vehicles, payments, and locations with completed Views for BI Analytics use at the end.

The project is built on Databricks (free tier) and emphasizes production-grade practices such as idempotency, cost optimization, and modular code. It showcases real-world data engineering skills applicable to cloud platforms like Azure, GCP, and AWS.

Tools and Technologies Used

- **PySpark:** For structured streaming, data ingestion, transformations, deduplication, and upserts.
- **Databricks:** Platform for compute, storage (Delta Lake), catalogs, schemas, volumes, and notebooks.
- **DBT (Data Build Tool):** For data modeling, incremental builds, snapshots, sources, and orchestration (using DBT Cloud free tier).
- **Jinja:** Templating engine within DBT for dynamic SQL generation (e.g., loops and conditionals).
- **Python:** For modular utilities, classes, and reusable transformation logic.
- **Delta Lake:** Storage format for reliable, ACID-compliant tables with merge capabilities.
- **YAML:** For DBT configurations, sources, and snapshots.
- **SQL:** Core language for DBT models and PySpark transformations.
- **Git:** Version control for DBT projects (integrated in DBT Cloud).
- **CSV Files:** Source data format for the pseudo-Uber dataset.

Relevant Skills Demonstrated

- PySpark Structured Streaming for incremental ingestion.
- Dynamic schema handling and notebook automation.
- Data cleaning techniques (e.g., regex, concatenation, conditional columns).
- Object-oriented programming in Python for reusable transformations.
- Deduplication using window functions and hashing.
- Upsert operations with Delta Tables and conditional merges.
- DBT modeling including incremental materializations, SCD Type 2 via snapshots, and lineage tracking.
- Jinja templating for dynamic SQL.
- Databricks environment management (catalogs, schemas, volumes, checkpoints).
- Data quality and governance (audit timestamps, idempotent processing).
- Modular and scalable pipeline design for production readiness.

Step-by-Step Project Build Process

Step 1: Environment Setup on Databricks

I started by setting up the environment in Databricks using the free community edition. This involved creating a serverless compute cluster with autoscaling to handle processing efficiently without manual VM management.

- Created a catalog named pyspark_dbt.
- Defined schemas: source for raw data, bronze for ingested data, silver for transformed data, and gold for modeled data.
- Set up a volume source_data under the source schema, with subdirectories for each entity (e.g., trips, customers, drivers, vehicles, payments, locations).
- Uploaded CSV files into the respective subdirectories.
- Created a checkpoint volume under the bronze schema to store metadata for incremental processing.

The screenshot shows the Databricks interface with the following details:

- Header:** Databricks Free Edition
- Search Bar:** Search d...
- New Button:** + New
- Sidebar:** Home, Workspace, Recents, Catalog (selected), Jobs & Pipelines, Compute, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie, Alerts, Query History, SQL Warehouses.
- Catalog View:** Catalog tab selected. Environment: Serverless Starter Warehouse, Serverless, 2XS.
 - Type to search... input field.
 - For you (button) and All (button) filters.
 - My organization
 - > workspace
 - > system
 - > pysparkdbt
 - > bronze
 - Tables (6)
 - customers
 - drivers
 - locations
 - payments
 - trips
 - vehicles
 - Volumes (1)
 - checkpoint

Databricks catalog and schemas setup

The screenshot shows the Databricks Catalog Explorer interface. At the top, there's a breadcrumb navigation: Catalog Explorer > pysparkdbt > source >. Below it is a header with a folder icon, the volume name "source_data", a refresh icon, a star icon, three dots for more options, a "Share" button, and a blue "Upload to this volume" button. Underneath the header, there are tabs for "Overview" (which is selected), "Files", "Details", and "Permissions".

Description

Two buttons are visible: "AI generate" (with a star icon) and "Add".

The main area shows a file tree: "/Volumes/pysparkdbt/source/source_data / customers". To the right of the tree are "Refresh" and "Create directory" buttons.

A search bar below the tree has the placeholder "Filter files and directori...".

A table lists the contents of the volume:

Name	Size	Last modified
customers.csv	20.10 KB	12 days ago

Source data volume with uploaded CSV files

Step 2: Bronze Layer - Incremental Data Ingestion with PySpark

In this step, I used PySpark Structured Streaming to ingest raw CSV data incrementally into Delta tables in the bronze layer. This ensures only new files are processed, making the pipeline efficient.

- Defined a list of entities: ["customers", "trips", "vehicles", "drivers", "payments", "locations"].
- In a Databricks notebook, used a loop to dynamically process each entity:
 - Read streaming data with `spark.readStream.format("csv").option("header", "true").option("inferSchema", "true")`.
 - Extracted schema dynamically from a batch DataFrame to avoid hardcoding.
 - Wrote to bronze tables using `format("delta").outputMode("append").option("checkpointLocation", "<path>").trigger(once=True).start()`.
- This created 6 Delta tables in the bronze schema, preserving raw data with checkpoints for fault tolerance.

The screenshot shows a Databricks notebook titled "bronze_ingestion". The notebook interface includes a header with File, Edit, View, Run, Help, Python, Tabs: ON, and a note that the last edit was 11 days ago. Below the header is a code editor window. The code is a PySpark script for a dynamic ingestion loop:

```

Dec 10, 2025          6
● for entity in entities:
    df_batch = spark.read.format("csv")\
        .option("header", True)\n        .option("inferSchema", True)\n        .load(f"/Volumes/pysparkdbt/source/source_data/{entity}/")
    schema_entity = df_batch.schema
    df = spark.readStream.format("csv")\
        .option("header", True)\n        .schema(schema_entity)\n        .load(f"/Volumes/pysparkdbt/source/source_data/{entity}")
    df.writeStream.format("delta")\
        .outputMode("append")\n        .option("checkpointLocation", f"/Volumes/pysparkdbt/bronze/checkpoint/{entity}")\n        .option("mergeSchema", "true")\n        .trigger(once=True)\n        .toTable(f"pysparkdbt.bronze.{entity}")

```

At the bottom of the code editor, there is a link "See performance (6)".

PySpark code for dynamic ingestion loop

The screenshot shows the Databricks catalog interface. On the left is a sidebar with navigation links: New, Home, Workspace, Recents, Catalog (which is selected), Jobs & Pipelines, Compute, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie, Alerts, and Query History. The main area is titled "Catalog" and shows the "Serverless Starter Warehouse" environment. It features a search bar with "Type to search..." and a filter icon. Below the search bar are tabs for "For you" and "All". The catalog tree structure is displayed under "pysparkdbt/bronze/Tables (6)":

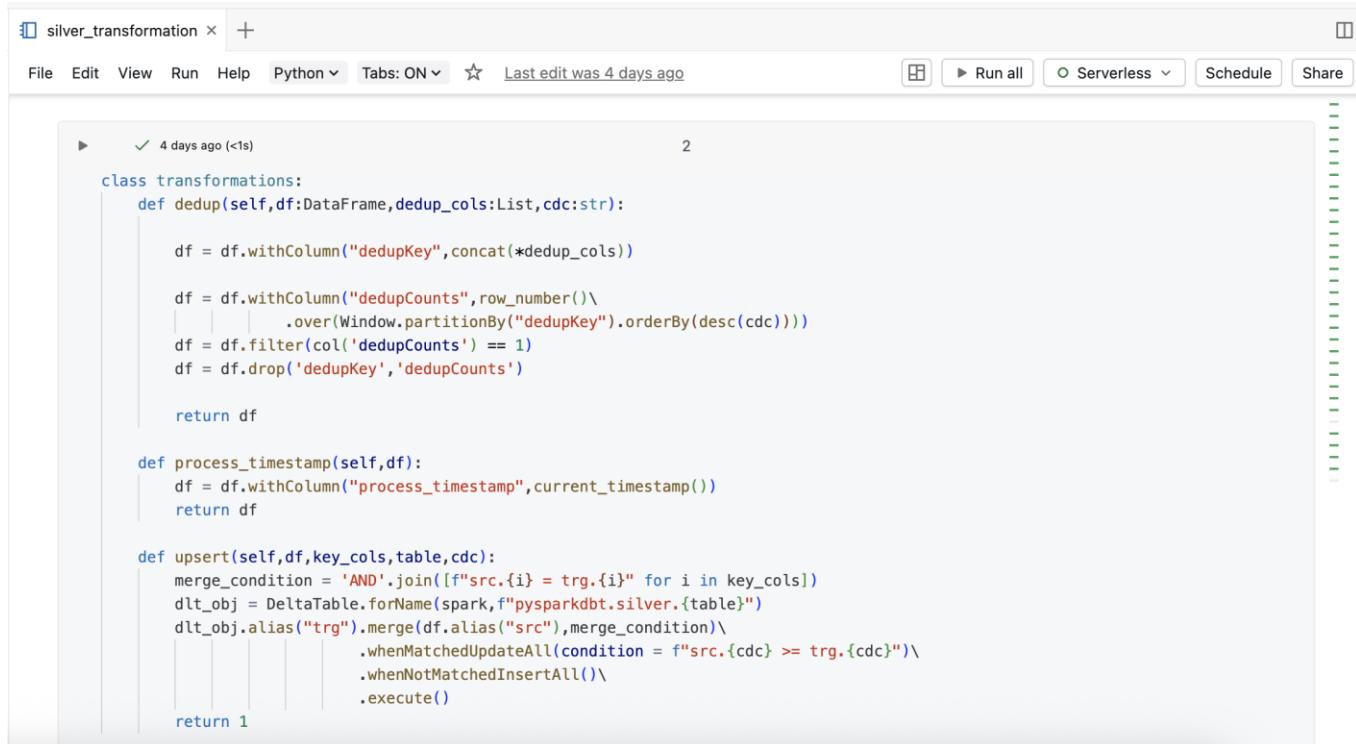
- customers
- drivers
- locations
- payments
- trips
- vehicles

Bronze layer Delta tables in Databricks catalog

Step 3: Silver Layer - Data Transformation with PySpark

Here, I focused on cleaning and enriching the data using modular Python code. I created a utility file `custom_utils.py` with a `Transformations` class to encapsulate reusable logic.

- Read data from bronze tables: `spark.read.table("bronze.<entity>")`.
- Applied generic transformations:
 - Deduplication: Used window functions (`row_number` over partitioned hash keys) to keep the latest records based on a change data capture (CDC) timestamp.
 - Added `process_timestamp` using `current_timestamp()` for auditing.
 - Upserts: Checked if silver table exists; if yes, performed Delta merge with dynamic conditions (e.g., update if source timestamp is newer).
- Entity-specific cleaning (examples):
 - Customers: Extracted email domains with `split`, cleaned phone numbers with `regexp_replace`, concatenated full names with `concat_ws`.
 - Payments: Added conditional `online_payment_status` using `when().otherwise()`.
 - Vehicles: Converted `make` to uppercase.
- Looped over entities to apply transformations and upsert into `silver.<entity>` tables.



The screenshot shows a Databricks notebook titled "silver_transformation". The notebook interface includes a header with tabs for File, Edit, View, Run, Help, Python (selected), Tabs: ON, Last edit was 4 days ago, and various execution and sharing buttons. The main code area contains the following Python code:

```
class Transformations:
    def dedup(self, df: DataFrame, dedup_cols: List, cdc: str):
        df = df.withColumn("dedupKey", concat(*dedup_cols))

        df = df.withColumn("dedupCounts", row_number()\
            .over(Window.partitionBy("dedupKey").orderBy(desc(cdc))))
        df = df.filter(col('dedupCounts') == 1)
        df = df.drop('dedupKey', 'dedupCounts')

        return df

    def process_timestamp(self, df):
        df = df.withColumn("process_timestamp", current_timestamp())
        return df

    def upsert(self, df, key_cols, table, cdc):
        merge_condition = 'AND'.join(['src.{i} = trg.{i}' for i in key_cols])
        dlt_obj = DeltaTable.forName(spark, f"pysparkdbt.silver.{table}")
        dlt_obj.alias("trg").merge(df.alias("src"), merge_condition)\n            .whenMatchedUpdateAll(condition = f"src.{cdc} >= trg.{cdc}")\n            .whenNotMatchedInsertAll()\n            .execute()

        return 1
```

Transformations class in custom_utils.py

The screenshot shows a Databricks notebook titled "silver_transformation". The notebook contains several code cells, each with a timestamp and a numerical identifier (6, 7, 8, 9, 10, 12) in the top right corner. The code performs various data transformations on a DataFrame named "df_cust".

```

6
df_cust = spark.read.table("pysparkdbt.bronze.customers")
> df_cust: pyspark.sql.connect.DataFrame = [customer_id: integer, first_name: string ... 6 more fields]

7
df_cust = df_cust.withColumn('domain', split(col('email'), '@')[1])
> df_cust: pyspark.sql.connect.DataFrame = [customer_id: integer, first_name: string ... 7 more fields]

8
df_cust = df_cust.withColumn("phone_number", regexp_replace("phone_number", r"^\d{10}", ""))
> df_cust: pyspark.sql.connect.DataFrame = [customer_id: integer, first_name: string ... 7 more fields]

9
df_cust = df_cust.withColumn("full_name", concat_ws(col('first_name'), col('last_name')))
df_cust = df_cust.drop('first_name', 'last_name')
> df_cust: pyspark.sql.connect.DataFrame = [customer_id: integer, email: string ... 6 more fields]

10
cust_obj = transformations()
cust_df_trns = cust_obj.dedup(df_cust[['customer_id']], 'last_updated_timestamp')
display(cust_df_trns)

```

A tooltip "Microsoft OneNote" appears over the "display" command in cell 10.

```

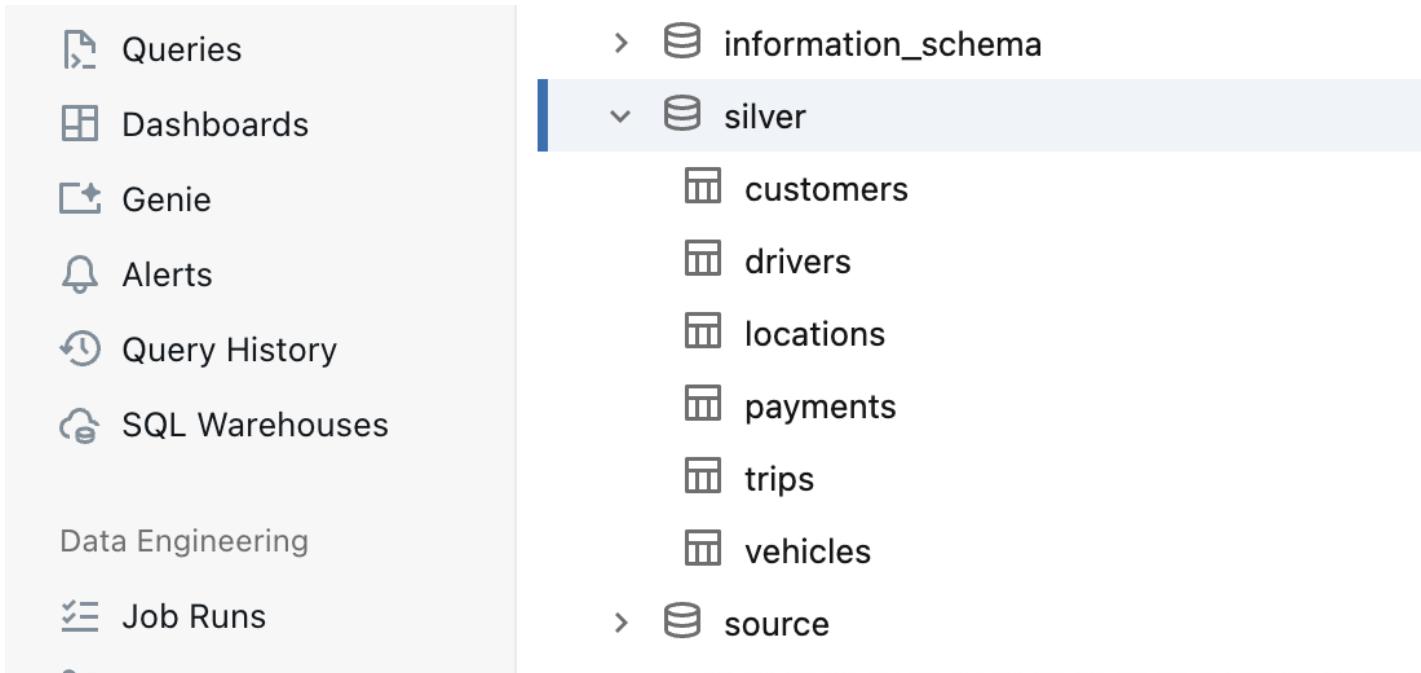
12
#this is the upsert function

if not spark.catalog.tableExists("pysparkdbt.silver.customers"):
    df_cust.write.format("delta")\
        .mode("append")\
        .saveAsTable("pysparkdbt.silver.customers")
else:
    cust_obj.upsert(df_cust, ['customer_id'], 'customers', 'last_updated_timestamp')

> See performance (2)

```

Example PySpark code for customer transformations and upsert



Silver layer tables after processing

Step 4: DBT Setup and Integration

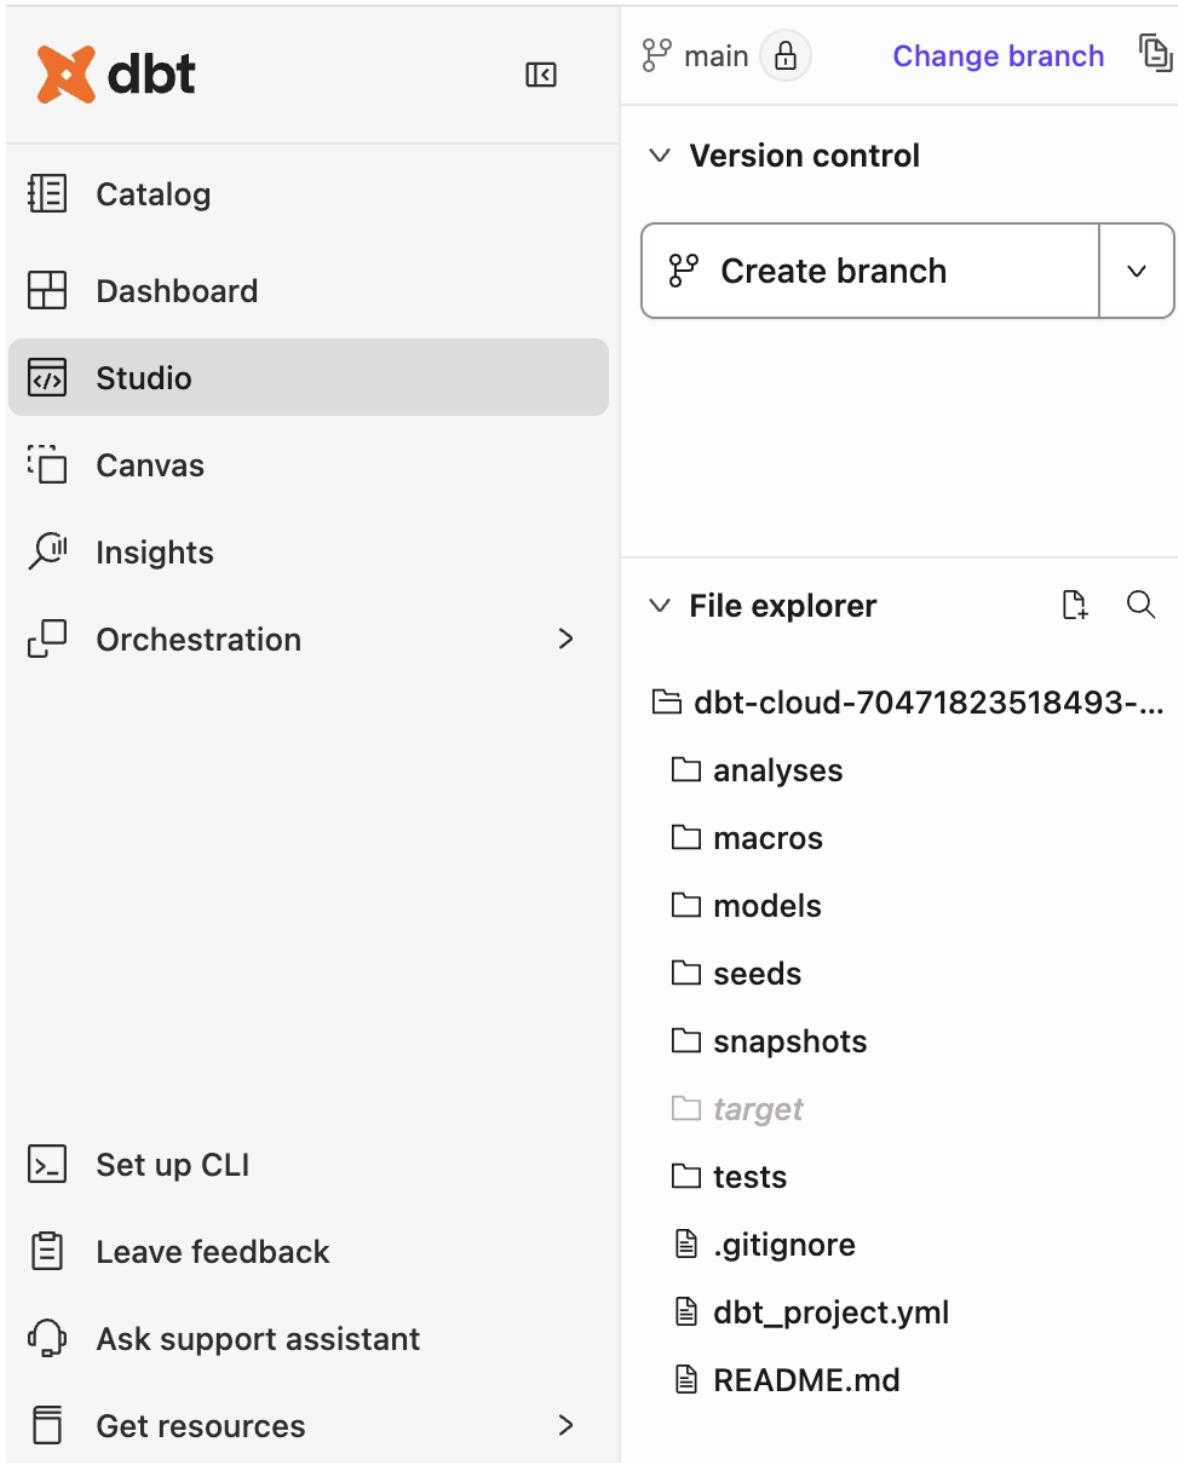
I shifted to DBT Cloud (free tier) for the modeling phase. Connected DBT to Databricks using hostname, HTTP path, access token, catalog, and schema details.

- Initialized a DBT project in DBT Cloud Studio.
- Created a development branch for safe experimentation.
- Defined project structure: models/silver/, models/gold/, sources/, snapshots/, macros/.

- Customized schema generation with a macro generate_schema_name.sql to use clean names like silver without prefixes.

The screenshot shows the DBT Cloud interface for setting up a connection to Databricks. On the left, there's a sidebar with various navigation options: Settings, Account, Projects, Integrations, Connections (which is highlighted in blue), Users, Billing, API tokens, Personal tokens, Service tokens, Webhooks, Your profile, Personal profile, Password & Security, Credentials, CLI, VS Code Extension, and Notification settings. The main area is titled "Databricks_Con" and shows "Connection settings". Under "Type", it's set to "Databricks". The "Connection name" is "Databricks_Con". In the "Settings" section, there's a note about IP access and a field for "Server Hostname" which is currently redacted. Below that is a placeholder text: "The hostname of the Databricks cluster or SQL warehouse." At the top right of the main area is an "Edit" button.

DBT Cloud connection setup to Databricks



DBT project structure in Studio

Step 5: Gold Layer - Dimensional Modeling with DBT

This step involved building a star schema with dimension and fact tables using DBT models, sources, and snapshots.

- Defined sources in sources.yml to register silver tables for lineage.
 - Created incremental models (e.g., silver/trips.sql):
 - Used config(materialized='incremental', unique_key='trip_id').
 - Applied Jinja for dynamic SELECT: Looped over column lists with {% for %}...{% if not loop.last %},{% endif %}.
 - Filtered incremental data with is_incremental() macro and last_updated_timestamp.
 - Implemented SCD Type 2 for dimensions via snapshots (e.g., dim_customers.yml):
 - Strategy: timestamp with unique_key and updated_at.
 - Added dbt_valid_from and dbt_valid_to for versioning.
 - Created fact table fact_trips referencing silver models with ref('trips') and joining dimensions.
 - Ran dbt run for models and dbt snapshot for SCDs.
 - Viewed lineage and documentation in DBT.

The screenshot shows the dbt Cloud interface. On the left, the sidebar has icons for version control, file explorer, and search. The 'File explorer' section is expanded, showing a tree structure with nodes for analyses, macros, models (gold, silver), sources, and seeds. The 'sources' node is selected, and its child 'sources.yml' file is also selected. The main content area shows the 'sources.yml' file's JSON content:

```
models > sources > sources.yml

dbt_yml_files (dbt_yml_files-latest.json)
version: 2

sources:
  - name: source_bronze
    database: pysparkdbt
    schema: bronze
    tables:
      - name: trips
        Generate model

  - name: source_silver
    database: pysparkdbt
    schema: silver
    tables:
      - name: customers
        Generate model
      - name: locations
        Generate model
      - name: drivers
        Generate model
      - name: payments
        Generate model
      - name: vehicles
```

```

models > silver > trips.sql

1  {{ config(
2    materialized ='incremental',
3    unique_key = 'trip_id'
4  )}}
5  }}
6
7
8
9  {% set cols = ['trip_id','vehicle_id','customer_id','driver_id','trip_start_time','trip_end_time','distance_km','fare_amount','last_updated_timestamp'] %}
10
11
12  SELECT
13    {% for col in cols %}
14      {{ col }}
15      {% if not loop.last %}
16          ,
17      {% endif %}
18    {% endfor %}
19  FROM
20  {{source('source_bronze','trips')}}
21  {% if is_incremental() %}
22  WHERE
23    last_updated_timestamp > (SELECT coalesce(MAX(last_updated_timestamp), '1900-01-01') FROM {{this}})
24  {% endif %}

```

Example DBT model with Jinja templating

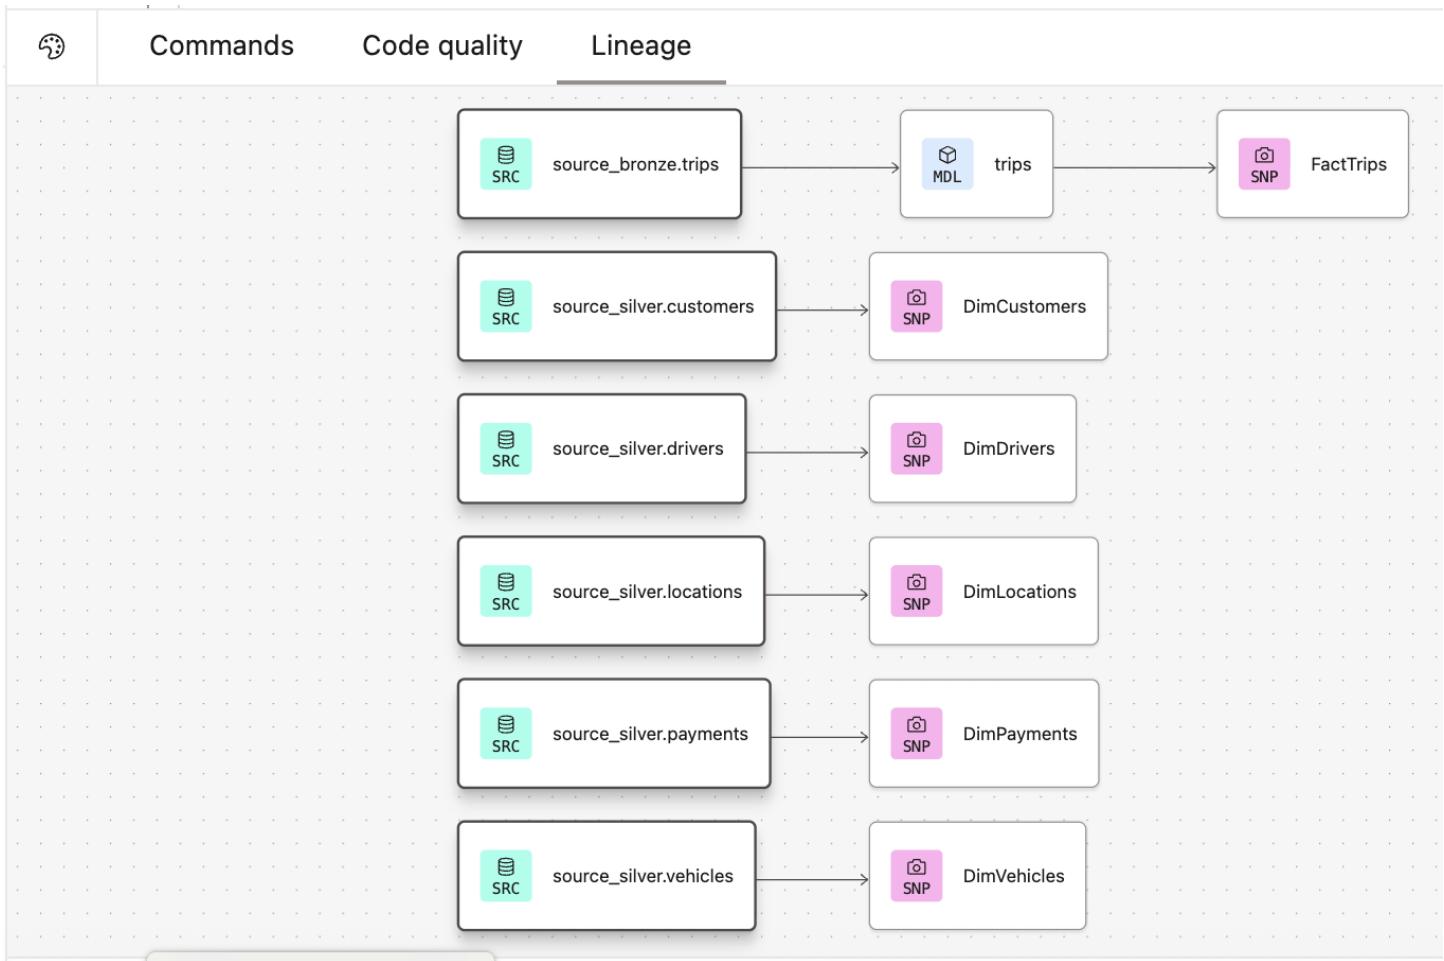
```

snapshots > SCDs.yml

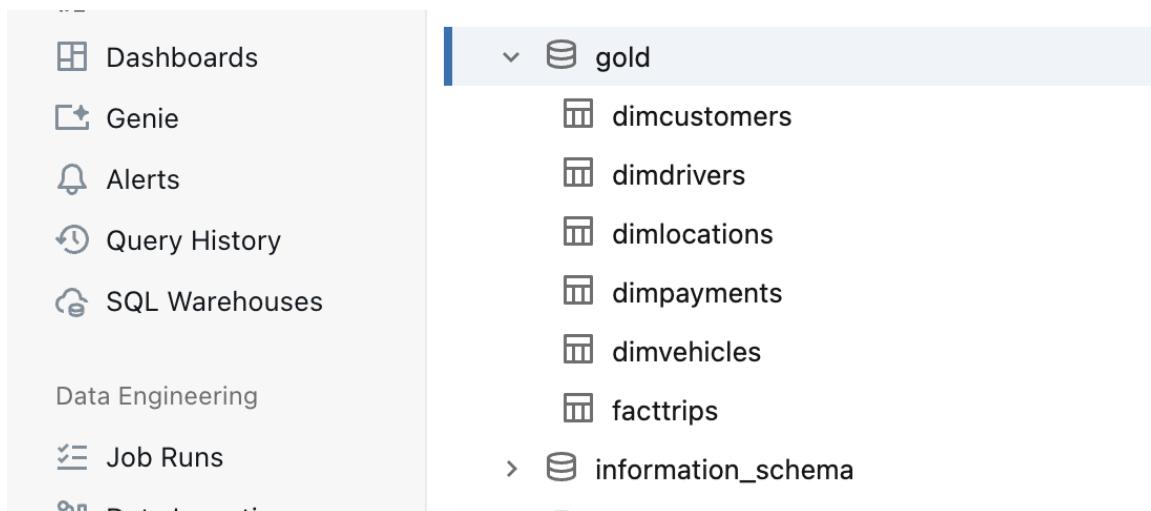
1  snapshots:
2    - name: DimCustomers
3      relation: source('source_silver', 'customers')
4      config:
5        schema: gold
6        database: pysparkdbt
7        unique_key: customer_id
8        strategy: timestamp
9        updated_at: last_updated_timestamp
10       dbt_valid_to_current: "to_date('9999-12-31')"
11
12    - name: DimLocations
13      relation: source('source_silver', 'locations')
14      config:
15        schema: gold
16        database: pysparkdbt
17        unique_key: location_id
18        strategy: timestamp
19        updated_at: last_updated_timestamp
20       dbt_valid_to_current: "to_date('9999-12-31')"
21
22    - name: DimDrivers
23      relation: source('source_silver', 'drivers')
24      config:
25        schema: gold
26        database: pysparkdbt
27        unique_key: driver_id
28        strategy: timestamp
29        updated_at: last_updated_timestamp
30       dbt_valid_to_current: "to_date('9999-12-31')"
31

```

DBT snapshot configuration for SCD



DBT lineage graph



Gold layer tables in Databricks (dimensions and facts)

Step 6: Final Deployment and Testing

- Committed changes to Git in DBT Cloud and merged to main.
- Ran full pipeline: Ingested new data in PySpark, then executed DBT jobs.
- Verified data in Databricks catalog and shared connection details for BI tools (e.g., Power BI).
- Ensured idempotency by re-running with no duplicates or errors.

> System logs

All 6 Pass 6 Warn 0 Error 0 Skip 0 Running 0

>	✓ DimCustomers	17.26s
>	✓ DimDrivers	16.03s
>	✓ DimLocations	15.90s
>	✓ DimPayments	16.45s
>	✓ DimVehicles	8.98s
>	✓ FactTrips	6.92s

DBT run output and success logs

Step 7: Created Analytical Views for Business Intelligence

- Developed 3 SQL views in Databricks gold schema to support BI and analytics.
 - Utilized **CTEs** for modular query structure and **window functions** (RANK(), ROW_NUMBER(), AVG() OVER) for rankings and rolling averages.
-
- Views answer key business questions:

1. vw_city_payment_performance: Regional payment success rates and 3-month rolling fare averages for operational insights.

```

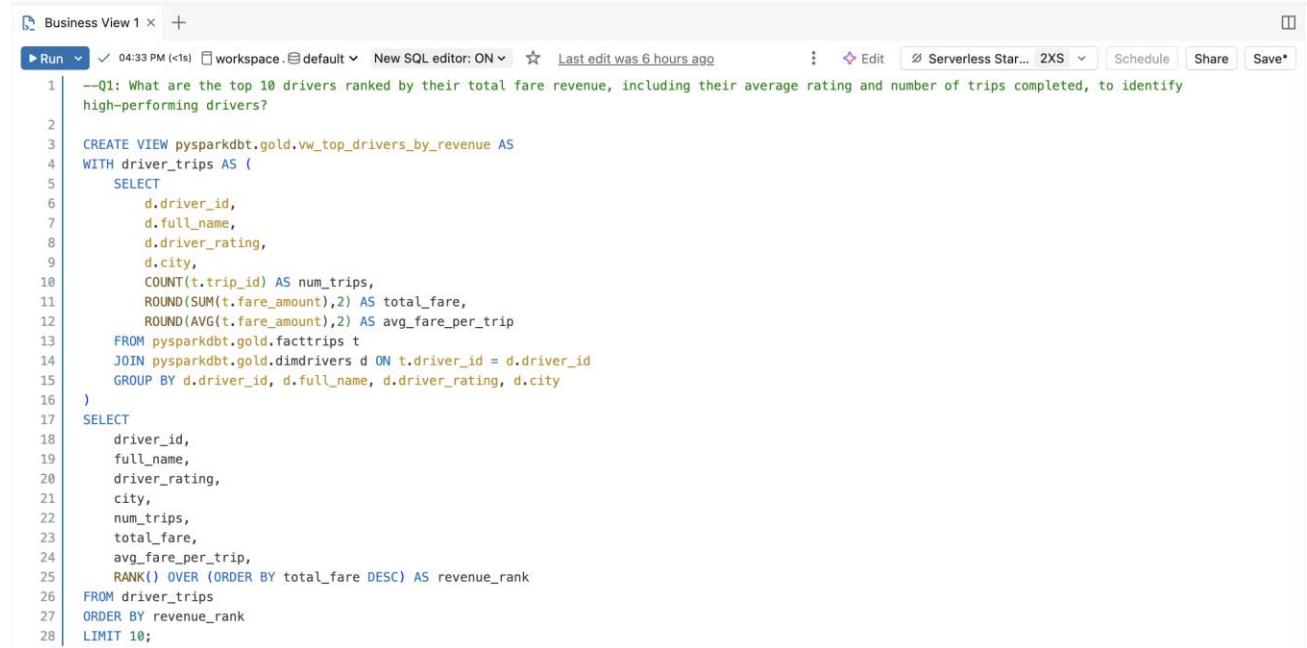
1 --Q3:What is the average fare amount and payment success rate per city, including a rolling average of fares over the last 3 months (assuming transaction_time in payments), to monitor regional performance?
2
3 CREATE VIEW pysparkdbt.gold.vw_city_payment_performance AS
4 WITH trip_payments AS (
5     SELECT
6         t.trip_id,
7         d.city AS driver_city, -- Using driver city as proxy for trip city
8         p.amount,
9         p.payment_status,
10        p.transaction_time,
11        CASE WHEN p.payment_status = 'success' THEN 1 ELSE 0 END AS success_flag
12    FROM pysparkdbt.gold.facttrips t
13    JOIN pysparkdbt.gold.dimpayments p ON t.trip_id = p.trip_id
14    JOIN pysparkdbt.gold.dimdrivers d ON t.driver_id = d.driver_id
15 ),
16 monthly_aggregates AS (
17     SELECT
18         driver_city AS city,
19         DATE_TRUNC('month', transaction_time) AS month,
20         ROUND(AVG(amount),2) AS avg_fare,
21         AVG(CAST(success_flag AS DOUBLE)) AS success_rate
22     FROM trip_payments
23     GROUP BY driver_city, DATE_TRUNC('month', transaction_time)
24 )
25 SELECT
26     city,
27     month,
28     avg_fare,
29     success_rate,
30     ROUND(AVG(avg_fare) OVER (PARTITION BY city ORDER BY month ROWS BETWEEN 2 PRECEDING AND CURRENT ROW), 2) AS rolling_3m_avg_fare
31 FROM monthly_aggregates
32 ORDER BY city, month DESC;

```

View 1 Query Results

	^{A_C} city	month	1.2 avg_fare	1.2 success_rate	1.2 rolling_3m_avg_fare
1	Brownburgh	2025-09-01T00:00:00.000+00:00	65.54	0	57.99
2	Brownburgh	2025-08-01T00:00:00.000+00:00	54.7	0	58.31
3	Brownburgh	2025-07-01T00:00:00.000+00:00	53.72	0	60.11
4	Brownburgh	2025-06-01T00:00:00.000+00:00	66.5	0	66.5
5	Charlesborough	2025-09-01T00:00:00.000+00:00	44.42	0	54.94
6	Charlesborough	2025-08-01T00:00:00.000+00:00	47.8	0	56.12
7	Charlesborough	2025-07-01T00:00:00.000+00:00	72.6	0	60.28
8	Charlesborough	2025-06-01T00:00:00.000+00:00	47.96	0	47.96
9	Christianshire	2025-09-01T00:00:00.000+00:00	63.49	0	56.32
10	Christianshire	2025-08-01T00:00:00.000+00:00	59.18	0	52.73

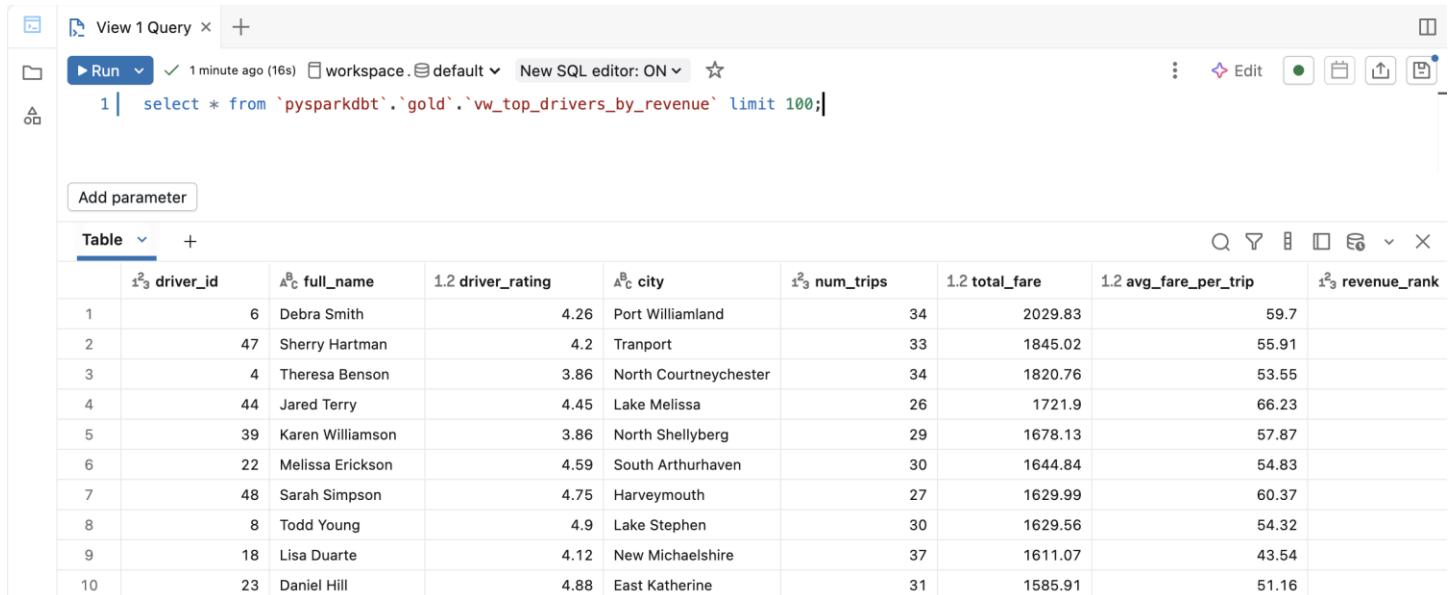
2. **vw_top_drivers_by_revenue:** Top 10 drivers by total fare, trip counts, and ratings for performance analysis.



```

1 --Q1: What are the top 10 drivers ranked by their total fare revenue, including their average rating and number of trips completed, to identify
2
3 CREATE VIEW pysparkdbt.gold.vw_top_drivers_by_revenue AS
4 WITH driver_trips AS (
5     SELECT
6         d.driver_id,
7         d.full_name,
8         d.driver_rating,
9         d.city,
10        COUNT(t.trip_id) AS num_trips,
11        ROUND(SUM(t.fare_amount),2) AS total_fare,
12        ROUND(AVG(t.fare_amount),2) AS avg_fare_per_trip
13    FROM pysparkdbt.gold.facttrips t
14    JOIN pysparkdbt.gold.dimdrivers d ON t.driver_id = d.driver_id
15    GROUP BY d.driver_id, d.full_name, d.driver_rating, d.city
16 )
17     SELECT
18         driver_id,
19         full_name,
20         driver_rating,
21         city,
22         num_trips,
23         total_fare,
24         avg_fare_per_trip,
25         RANK() OVER (ORDER BY total_fare DESC) AS revenue_rank
26     FROM driver_trips
27     ORDER BY revenue_rank
28     LIMIT 10;
--
```

View 2 Query Results



View 1 Query

```

1 | select * from `pysparkdbt`.`gold`.`vw_top_drivers_by_revenue` limit 100;
```

Add parameter

#	driver_id	full_name	driver_rating	city	num_trips	total_fare	avg_fare_per_trip	revenue_rank
1	6	Debra Smith	4.26	Port Williamland	34	2029.83	59.7	
2	47	Sherry Hartman	4.2	Tranport	33	1845.02	55.91	
3	4	Theresa Benson	3.86	North Courtneychester	34	1820.76	53.55	
4	44	Jared Terry	4.45	Lake Melissa	26	1721.9	66.23	
5	39	Karen Williamson	3.86	North Shellyberg	29	1678.13	57.87	
6	22	Melissa Erickson	4.59	South Arthurhaven	30	1644.84	54.83	
7	48	Sarah Simpson	4.75	Harveymouth	27	1629.99	60.37	
8	8	Todd Young	4.9	Lake Stephen	30	1629.56	54.32	
9	18	Lisa Duarte	4.12	New Michaelshire	37	1611.07	43.54	
10	23	Daniel Hill	4.88	East Katherine	31	1585.91	51.16	

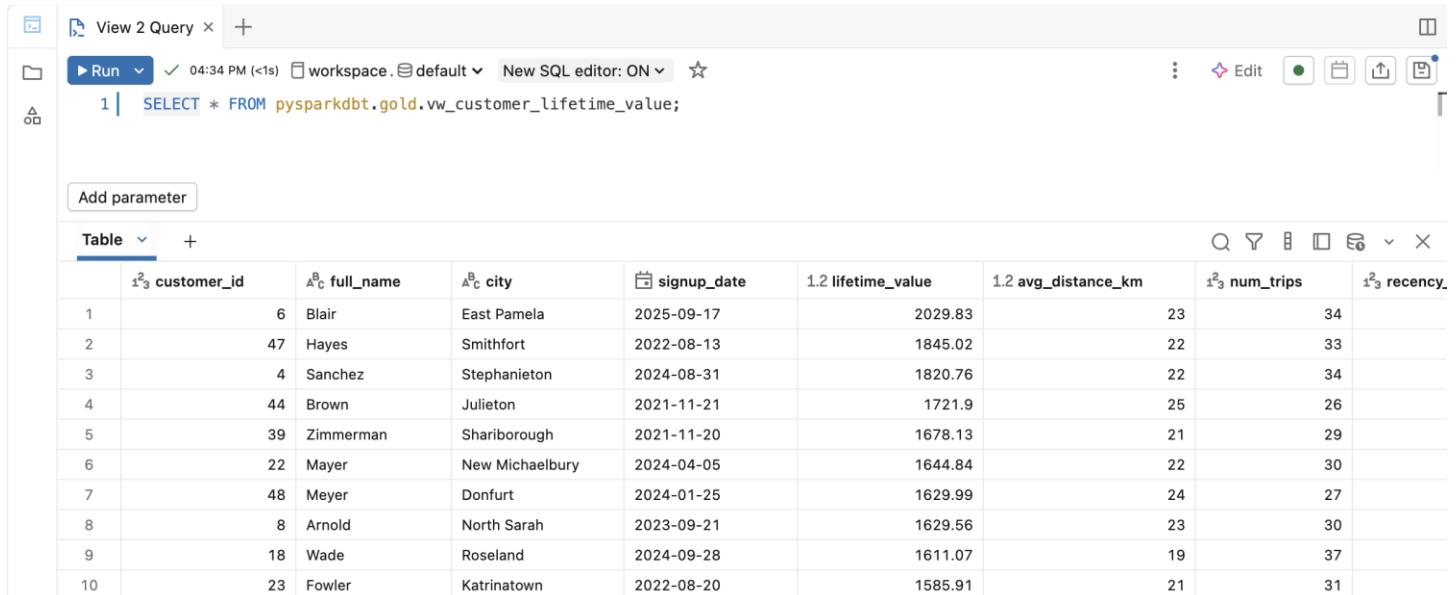
3. vw_customer_lifetime_value: Customer LTV, trip frequency, and recency for segmentation and retention strategies.



```

1 --Q2: For each customer, what is their lifetime value (total fare paid), average trip distance, and recency rank based on their last signup or trip date, to
2 analyze customer engagement?
3
4 CREATE VIEW pysparkdbt.gold.vw_customer_lifetime_value AS
5 WITH customer_stats AS (
6   SELECT
7     c.customer_id,
8     c.full_name,
9     c.city,
10    c.signup_date,
11    ROUND(SUM(t.fare_amount),2) AS lifetime_value,
12    ROUND(AVG(t.distance_km),0) AS avg_distance_km,
13    COUNT(t.trip_id) AS num_trips,
14    MAX(t.trip_id) AS last_trip_id -- Assuming trip_id increases over time as a proxy for recency
15  FROM pysparkdbt.gold.dimcustomers c
16  LEFT JOIN pysparkdbt.gold.facttrips t ON c.customer_id = t.customer_id
17  GROUP BY c.customer_id, c.full_name, c.city, c.signup_date
18 )
19 SELECT
20   customer_id,
21   full_name,
22   city,
23   signup_date,
24   lifetime_value,
25   avg_distance_km,
26   num_trips,
27   ROW_NUMBER() OVER (ORDER BY last_trip_id DESC) AS recency_rank
28 FROM customer_stats
29 ORDER BY lifetime_value DESC;
--
```

View 3 Query Results



View 2 Query

```

1 | SELECT * FROM pysparkdbt.gold.vw_customer_lifetime_value;
```

Add parameter

Table

#	customer_id	full_name	city	signup_date	lifetime_value	avg_distance_km	num_trips	recency
1	6	Blair	East Pamela	2025-09-17	2029.83		23	34
2	47	Hayes	Smithfort	2022-08-13	1845.02		22	33
3	4	Sanchez	Stephanieton	2024-08-31	1820.76		22	34
4	44	Brown	Julieton	2021-11-21	1721.9		25	26
5	39	Zimmerman	Shariborough	2021-11-20	1678.13		21	29
6	22	Mayer	New Michaelbury	2024-04-05	1644.84		22	30
7	48	Meyer	Donfurt	2024-01-25	1629.99		24	27
8	8	Arnold	North Sarah	2023-09-21	1629.56		23	30
9	18	Wade	Roseland	2024-09-28	1611.07		19	37
10	23	Fowler	Katrinatown	2022-08-20	1585.91		21	31

- Views query the star schema models, providing analysts clean, aggregated data without complex joins.

Project Summary

This end-to-end data engineering project transforms raw pseudo-Uber CSV data into a fully modeled analytics pipeline using PySpark for ingestion and transformation, and DBT for advanced modeling and orchestration on Databricks. It demonstrates proficiency in handling incremental loads, data quality, SCD Type 2, and dynamic SQL with Jinja, resulting in a scalable, production-ready data warehouse. Key outcomes include 6 bronze tables, cleaned silver layers with upserts, and a gold star schema with historical tracking. This portfolio project highlights my ability to build efficient, modular data pipelines, making it ideal for data engineering roles requiring expertise in big data tools, cloud platforms, and modern data stacks. The entire codebase is available in the repository for review.